# MIT 8103 Advanced Database Systems

Banking accounts and transactions · Isaac Adeiza · Student ID 301906391

## Case study

My case study is Banking accounts and transactions. I designed the database for bank staff reviewing customer accounts and transaction activity and students analyzing database design and query behavior, with separate tables named customer, account and bank_transaction. It supports these operations: open customer accounts and record account details; post deposits and withdrawals while maintaining each account balance; review transaction activity and summarize posted amounts; benchmark filtering transactions by status with and without a B-tree index using 25000 synthetic rows; compare the planned and observed query behavior without treating the discussion as a measured result.

## Portfolio 1: Database Design and Modelling

I organised the relational design into 3 tables. Primary keys identify records, foreign keys preserve the links, and checks enforce the stated business rules. An account balance must be zero or greater, preventing overdrafts. The normalisation examples show the sample data at 1NF, 2NF and 3NF. All 9 deliberately invalid inserts were rejected without changing the valid sample rows.

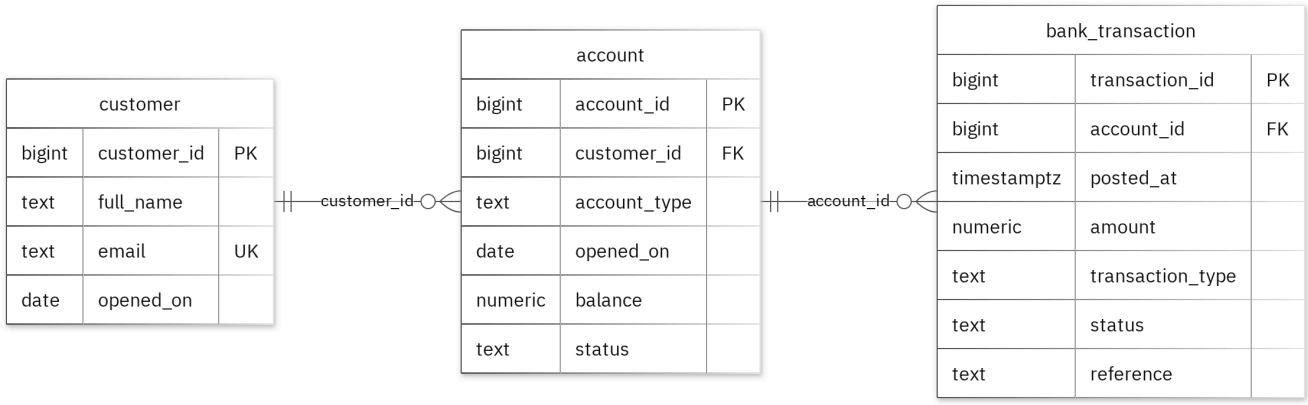



Figure 1. Tables, keys and relationships. PK: primary key; FK: foreign key; UK: unique value.

## Portfolio 2: Query Processing and Optimisation

I ran three SQL queries against the sample records. The join returned 2 rows and included Avery Morgan with checking. For the index comparison, I used 25,000 separate performance rows. The median lookup time changed from 1.482 ms to 0.17 ms. Both phases returned the same records. The measured plan used the new index. These timings apply to this lookup and environment.

## Portfolio 3: Transactions and Concurrency

I used a temporary counter on account to test COMMIT, ROLLBACK and READ COMMITTED isolation. The uncommitted update remained hidden from the other session, and ROLLBACK restored the earlier committed value. Without locking, two increments left the counter at 2. With SELECT FOR UPDATE, Session B waited for Session A and the final value was 3. The helper is confined to this exercise; the case's original business fields stay unchanged.

## Portfolio 4: NoSQL and Advanced Data Models

I used account_events for the case's event documents. MongoDB inserted 4 sample events and returned the expected results for 4 queries. Validation rejected an invalid date type and a duplicate event ID. Embedded snapshots keep related details with an event, while the details field can differ by event type. These are fictional events; PostgreSQL continues to enforce its relational rules.

## Repository and evidence

https://github.com/Iziik/Adeiza_Isaac_301906391

The portfolio folders contain the SQL and JavaScript, explanations, query plans, result tables and screenshots.

## Portfolio 5: Distributed and Cloud Database Exercise

I proposed three MongoDB replica members across separate failure zones. The discussion covers primary elections, majority acknowledgement, secondary lag and sharding trade-offs for account_events. Replication supports availability, while backups address recovery. PostgreSQL needs its own availability plan. This is the permitted discussion option; no live replica set or failover test was deployed.

## Individual reflection

I kept customer, account, bank_transaction in separate tables so their details could be maintained without repeating them in every related record. An account balance must be zero or greater, preventing overdrafts. The constraint tests check this alongside the required values and record links. One issue illustrated by the two-session exercise is that both sessions can read the same value before either update is committed. Locking the row before reading preserves both increments. The index comparison also shows why I need to check the execution plan as well as the final answer: the same result can require different amounts of database work.

## AI-use declaration

I used GitHub Copilot and the portfolio generator to help prepare the database model, scripts and explanations. The generator ran the PostgreSQL and MongoDB exercises, checked their outputs and captured the result images. I remain responsible for checking and explaining the submitted work.

## References


- PostgreSQL: constraints
- PostgreSQL: EXPLAIN
- PostgreSQL: isolation
- PostgreSQL: row locks
- MongoDB: data modelling
- MongoDB: replication
- MongoDB: sharding
- Mermaid: ER diagrams